

# *Lecture 16: Continued optimizations on GEMM kernels - Assembly-language level optimizations (via intrinsics)*

*Thursday March 23rd 2023*

# Logistics

- Homework #3 to be by tomorrow
- Due Friday March 31st
- Please check that you can use MKL (if you haven't already!). See the discussion in the March 2nd lecture, too.

# Today's lecture

- Continued focus on GEMM operations
- We will look into last-mile optimizations, that will involve explicitly looking into and writing assembly-level code.
- We will use intrinsics to access assembly-level functionality within the framework of the C compiler.

# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

RECAP

eg. no alias is bA failing to be treated as a constant (if it is one) because the compiler pessimistically thinks that it could be modified within the iteration. fnoalias essentially guarantees race free from the programmer

```
#include "MatMatMultiplyBlockHelper.h"

void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])
{
    for (int ii = 0; ii < BLOCK_SIZE; ii++)
        for (int kk = 0; kk < BLOCK_SIZE; kk++) {
#pragma omp simd
            for (int jj = 0; jj < BLOCK_SIZE; jj++)
                bC[ii][jj] += bA[ii][kk] * bB[kk][jj];
        }
}
```

*Instruct the compiler that no aliases are present*

```
icc -S MatMatMultiplyBlockHelper.cpp -qopenmp -xCOMMON-AVX512
    -fno-alias -o MatMatMultiplyBlockHelper.AVX512.NoAliases.s
```

# Assembly code

MatMatMultiplyBlockHelper.AVX512.NoAliases.s

```
# MatMatMultiplyBlockHelper(const float (&)[64][64], const float (&)[64][64], float (&)[64][64])
_Z25MatMatMultiplyBlockHelperRA64_A64_KfS2_RA64_A64_f:
# parameter 1: %rdi
# parameter 2: %rsi
# parameter 3: %rdx
[...]
..B1.3:                # Preds ..B1.3 ..B1.2
                        # Execution count [4.10e+03]
    vbroadcastss (%r9,%r10,4), %zmm4                #10.27
    incq         %r10                                #7.9
    vfmadd231ps  64(%r8,%rsi), %zmm4, %zmm0          #10.13
    vfmadd231ps  128(%r8,%rsi), %zmm4, %zmm1         #10.13
    vfmadd231ps  192(%r8,%rsi), %zmm4, %zmm3         #10.13
    vfmadd231ps  (%r8,%rsi), %zmm4, %zmm2            #10.13
    vmovups      %zmm0, 64(%rax,%rdx)                #10.13
    vmovups      %zmm1, 128(%rax,%rdx)               #10.13
    addq         $256, %r8                            #7.9
    cmpq         $64, %r10                            #7.9
    jb           ..B1.3                               # Prob 98%      #7.9
..B1.4:                # Preds ..B1.3
                        # Execution count [6.40e+01]
    incb         %cl                                    #6.5
    vmovups      %zmm3, 192(%rax,%rdx)               #10.13
    vmovups      %zmm2, (%rax,%rdx)                  #10.13
    addq         $256, %rax                            #6.5
    cmpb         $64, %cl                             #6.5
    jb           ..B1.2                               # Prob 98%      #6.5
```

# Assembly code

MatMatMultiplyBlockHelper.AVX512.NoAliases.s

RECAP

```
# MatMatMultiplyBlockHelper(const float (&)[64][64], const float (&)[64][64], float (&)[64][64])
_Z25MatMatMultiplyBlockHelperRA64_A64_KfS2_RA64_A64_f:
```

```
# parameter 1: %rdi
```

```
# parameter 2: %rsi
```

```
# parameter 3: %rdx
```

```
[...]
```

```
..B1.3:                                # Preds ..B1.3 ..B1.2
                                         # Execution count [4.10e+03]
```

```
vbroadcastss (%r9,%r10,4), %zmm4
```

```
incq        %r10
```

```
vfmadd231ps 64(%r8,%rsi), %zmm4, %zmm0
```

```
vfmadd231ps 128(%r8,%rsi), %zmm4, %zmm1
```

```
vfmadd231ps 192(%r8,%rsi), %zmm4, %zmm3
```

```
vfmadd231ps (%r8,%rsi), %zmm4, %zmm2
```

```
vmovups     %zmm0, 64(%rax,%rdx)
```

```
vmovups     %zmm1, 128(%rax,%rdx)
```

```
addq        $256, %r8
```

```
cmpq        $64, %r10
```

```
jb          ..B1.3                # Prob 98%
```

```
..B1.4:                                # Preds ..B1.3
                                         # Execution count [6.40e+01]
```

```
incb        %cl
```

```
vmovups     %zmm3, 192(%rax,%rdx)
```

```
vmovups     %zmm2, (%rax,%rdx)
```

```
addq        $256, %rax
```

```
cmpb        $64, %cl
```

```
jb          ..B1.2                # Prob 98%
```

#10.27

#7

#1 *Read bB[kk][jj] just once!*

#10.13

#10.13

#10.13

#10.13

#10.13

#7.9

#7.9

#7.9

#6.5

#10.13

#10.13

#6.5

#6.5

#6.5



# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

```
#include "MatMatMultiplyBlockHelper.h"

void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])
{
    for (int ii = 0; ii < BLOCK_SIZE; ii++)
        for (int kk = 0; kk < BLOCK_SIZE; kk++) {
#pragma omp simd
            for (int jj = 0; jj < BLOCK_SIZE; jj++)
                bC[ii][jj] += bA[ii][kk] * bB[kk][jj];
        }
}
```

# Assembly code

MatMatMultiplyBlockHelper.AVX512.NoAliases.s

```
# MatMatMultiplyBlockHelper(const float (&)[64][64], const float (&)[64][64], float (&)[64][64])
_Z25MatMatMultiplyBlockHelperRA64_A64_KfS2_RA64_A64_f:
```

```
# parameter 1: %rdi
```

```
# parameter 2: %rsi
```

```
# parameter 3: %rdx
```

```
[...]
```

```
..B1.3:                                # Preds ..B1.3 ..B1.2
                                         # Execution count [4.10e+03]
```

```
vbroadcastss (%r9,%r10,4), %zmm4
```

```
incq        %r10
```

```
vfmadd231ps 64(%r8,%rsi), %zmm4, %zmm0
```

```
vfmadd231ps 128(%r8,%rsi), %zmm4, %zmm1
```

```
vfmadd231ps 192(%r8,%rsi), %zmm4, %zmm3
```

```
vfmadd231ps (%r8,%rsi), %zmm4, %zmm2
```

```
vmovups     %zmm0, 64(%rax,%rdx)
```

```
vmovups     %zmm1, 128(%rax,%rdx)
```

```
addq        $256, %r8
```

```
cmpq        $64, %r10
```

```
jb          ..B1.3                # Prob 98%
```

```
..B1.4:                                # Preds ..B1.3
                                         # Execution count [6.40e+01]
```

```
incb        %cl
```

```
vmovups     %zmm3, 192(%rax,%rdx)
```

```
vmovups     %zmm2, (%rax,%rdx)
```

```
addq        $256, %rax
```

```
cmpb        $64, %cl
```

```
jb          ..B1.2                # Prob 98%
```

RECAP

*Do the multiplication & addition  
for the 64 entries of the row  
in 4 tightly-packed  
fused-multiply-add instructions  
(16 entries at a time)*

#10.13

#7.9

#7.9

#7.9

#6.5

#10.13

#10.13

#6.5

#6.5

#6.5



# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

**RECAP**

```
#include "MatMatMultiplyBlockHelper.h"

void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])
{
    for (int ii = 0; ii < BLOCK_SIZE; ii++)
        for (int kk = 0; kk < BLOCK_SIZE; kk++) {
#pragma omp simd
            for (int jj = 0; jj < BLOCK_SIZE; jj++)
                bC[ii][jj] += bA[ii][kk] * bB[kk][jj];
        }
}
```

# Assembly code

MatMatMultiplyBlockHelper.AVX512.NoAliases.s

RECAP

```
# MatMatMultiplyBlockHelper(const float (&)[64][64], const float (&)[64][64], float (&)[64][64])
_Z25MatMatMultiplyBlockHelperRA64_A64_KfS2_RA64_A64_f:
```

```
# parameter 1: %rdi
```

```
# parameter 2: %rsi
```

```
# parameter 3: %rdx
```

```
[...]
```

```
..B1.3:                                # Preds ..B1.3 ..B1.2
                                      # Execution count [4.10e+01]
    vbroadcastss (%r9,%r10,4), %zmm4
    incq        %r10
    vfmadd231ps 64(%r8,%rsi), %zmm4, %zmm0
    vfmadd231ps 128(%r8,%rsi), %zmm4, %zmm1
    vfmadd231ps 192(%r8,%rsi), %zmm4, %zmm2
    vfmadd231ps (%r8,%rsi), %zmm4, %zmm2
    vmovups     %zmm0, 64(%rax,%rdx)
    vmovups     %zmm1, 128(%rax,%rdx)
    addq        $256, %r8
    cmpq        $64, %r10
    jb          ..B1.3                # Prob 98%
..B1.4:                                # Preds ..B1.3
                                      # Execution count [6.40e+01]
    incb        %cl
    vmovups     %zmm3, 192(%rax,%rdx)
    vmovups     %zmm2, (%rax,%rdx)
    addq        $256, %rax
    cmpb        $64, %cl
    jb          ..B1.2                # Prob 98%
```

*Write the result back into  $bC[ii][jj]$   
using 4x 16-wide store (“move”) instructions  
(the compiler does a bit extra loop reordering,  
hence the split of the “move” instructions)*

src, dest

#10.13  
#10.13  
#7.9  
#7.9  
#7.9

#6.5  
#10.13  
#10.13  
#6.5  
#6.5  
#6.5

# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

RECAP

```
#include "MatMatMultiplyBlockHelper.h"

void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])
{
    for (int ii = 0; ii < BLOCK_SIZE; ii++)
        for (int kk = 0; kk < BLOCK_SIZE; kk++) {
#pragma omp simd
            for (int jj = 0; jj < BLOCK_SIZE; jj++)
                bC[ii][jj] += bA[ii][kk] * bB[kk][jj];
        }
}
```

*Building the entire executable*

*(it's important to only use the "no aliases" option on the isolated Block-matrix-multiply code file, as we do employ aliases widely elsewhere in the code, i.e. all the matrix casts ...)*

```
icc -c MatMatMultiplyBlockHelper.cpp -qopenmp -xCOMMON-AVX512 -fno-alias
```

```
icc main.cpp MatMatMultiply.cpp Utilities.cpp MatMatMultiplyBlockHelper.o
    -xCOMMON-AVX512 -mkl -xHost
```

# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

RECAP

```
#include "MatMatMultiplyBlockHelper.h"
```

```
void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])
```

```
{
    for (int ii = 0; ii < BLOCK_SIZE; ii++)
        for (int kk = 0; kk < BLOCK_SIZE; kk++) {
```

```
#pragma omp simd
```

```
    for (int jjRunning candidate kernel for correctness test ... [Elapsed time : 41.0696ms]
        bC[ii][Running reference kernel for correctness test ... [Elapsed time : 39.3287ms]
```

```
    }
    Discrepancy between two methods : 0.000160217
```

```
    Running kernel for performance run # 1 ... [Elapsed time : 23.0806ms]
```

```
    Running kernel for performance run # 2 ... [Elapsed time : 22.582ms]
```

```
    Running kernel for performance run # 3 ... [Elapsed time : 22.1493ms]
```

```
    Running kernel for performance run # 4 ... [Elapsed time : 21.97ms]
```

```
    Running kernel for performance run # 5 ... [Elapsed time : 22.8829ms]
```

```
    Running kernel for performance run # 6 ... [Elapsed time : 22.8694ms]
```

```
    Running kernel for performance run # 7 ... [Elapsed time : 21.5426ms]
```

```
    Running kernel for performance run # 8 ... [Elapsed time : 22.7655ms]
```

```
    [...]
```

*1.8x the runtime of MKL code!*

## Execution:

## Last-mile optimizations ...

*So far, we looked at instances where performance of transformed (or refactored) code would be best understood by looking at assembly code.*

*However, we did not write such assembly code directly ... today we will.*

*These optimizations can be CPU-specific and compiler-specific (we will focus on machines that support AVX2 and/or AVX512)*

*You can look up what your processor supports at [ark.intel.com](http://ark.intel.com) (if Intel CPU)*

*Goal : Understand the nature, style and motivation of these optimizations  
You will not be required to reproduce such optimizations in homework or exams*

## Note on examples

- Makefiles for Linux (should be ok in OS X too) are included
- Typing “make assembly” should produce the assembly code for `MatMatMultiplyBlockHelper.cpp` (with an `.s` extension)
- Many directories will include the assembly file into the repository, for reference
  - All examples in **GEMM\_Test\_1\_XXX**



# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

```
#include "MatMatMultiplyBlockHelper.h"

void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])
{
    for (int ii = 0; ii < BLOCK_SIZE; ii++)
        for (int kk = 0; kk < BLOCK_SIZE; kk++) {
#pragma omp simd
            for (int jj = 0; jj < BLOCK_SIZE; jj++)
                bC[ii][jj] += bA[ii][kk] * bB[kk][jj];
        }
}
```

*Build with:*

```
icc -c MatMatMultiplyBlockHelper.cpp -qopenmp -xCOMMON-AVX512 -fno-alias
icc main.cpp MatMatMultiply.cpp Utilities.cpp MatMatMultiplyBlockHelper.o
    -xCOMMON-AVX512 -mkl -xHost
```

# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

```
#include "MatMatMultiplyBlockHelper.h"
```

```
void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])
```

```
{
    for (int ii = 0; ii < BLOCK_SIZE; ii++)
        for (int kk = 0; kk < BLOCK_SIZE; kk++) {
```

```
#pragma omp simd
```

```
    for (int jjRunning candidate kernel for correctness test ... [Elapsed time : 41.0696ms]
        bC[ii][Running reference kernel for correctness test ... [Elapsed time : 39.3287ms]
```

```
    }
    Discrepancy between two methods : 0.000160217
```

```
    Running kernel for performance run # 1 ... [Elapsed time : 23.0806ms]
```

```
    Running kernel for performance run # 2 ... [Elapsed time : 22.582ms]
```

```
    Running kernel for performance run # 3 ... [Elapsed time : 22.1493ms]
```

```
    Running kernel for performance run # 4 ... [Elapsed time : 21.97ms]
```

```
    Running kernel for performance run # 5 ... [Elapsed time : 22.8829ms]
```

```
    Running kernel for performance run # 6 ... [Elapsed time : 22.8694ms]
```

```
    Running kernel for performance run # 7 ... [Elapsed time : 21.5426ms]
```

```
    Running kernel for performance run # 8 ... [Elapsed time : 22.7655ms]
```

```
    [...]
```

*1.8x the runtime of MKL code!*

## Execution:

# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

```
#include "MatMatMultiplyBlockHelper.h"

void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])
{
    for (int ii = 0; ii < BLOCK_SIZE; ii++)
        for (int kk = 0; kk < BLOCK_SIZE; kk++) {
#pragma omp simd
            for (int jj = 0; jj < BLOCK_SIZE; jj++)
                bC[ii][jj] += bA[ii][kk] * bB[kk][jj];
        }
}
```

*We still had to rely on OpenMP to (hopefully) generate SIMD code*

# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

*DenseAlgebra/GEMM\_Test\_1\_0\_avx512*

```
#include "MatMatMultiplyBlockHelper.h"
#include "immintrin.h"

void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])
{
    static constexpr int nW = 16; // Width of SIMD vectors

    for (int ii = 0; ii < BLOCK_SIZE; ii++)
        for (int kk = 0; kk < BLOCK_SIZE; kk++)
            for (int jj = 0; jj < BLOCK_SIZE; jj += nW) {
                __m512 vB = _mm512_load_ps(&bB[kk][jj]);
                __m512 vA = _mm512_set1_ps(bA[ii][kk]);
                __m512 vC = _mm512_load_ps(&bC[ii][jj]);
                vC = _mm512_fmadd_ps(vA, vB, vC);
                _mm512_store_ps(&bC[ii][jj], vC);
            }
}
```

*This code explicitly intended for an AVX512-compatible CPU ...*

# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

*DenseAlgebra/GEMM\_Test\_1\_0\_avx2*

```
#include "MatMatMultiplyBlockHelper.h"
#include "immintrin.h"

void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])
{
    static constexpr int nW = 8; // Width of SIMD vectors

    for (int ii = 0; ii < BLOCK_SIZE; ii++)
        for (int kk = 0; kk < BLOCK_SIZE; kk++)
            for (int jj = 0; jj < BLOCK_SIZE; jj += nW) {
                __m256 vB = _mm256_load_ps(&bB[kk][jj]);
                __m256 vA = _mm256_set1_ps(bA[ii][kk]);
                __m256 vC = _mm256_load_ps(&bC[ii][jj]);
                vC = _mm256_fmadd_ps(vA, vB, vC);
                _mm256_store_ps(&bC[ii][jj], vC);
            }
}
```

*... but a version is provided for AVX2-compatible CPUs*

# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

*DenseAlgebra/GEMM\_Test\_1\_0\_avx512*

```
#include "MatMatMultiplyBlockHelper.h"  
#include "immintrin.h"
```

*We will use assembly intrinsics ...*

```
void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],  
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])  
{  
    static constexpr int nW = 16; // Width of SIMD vectors  
  
    for (int ii = 0; ii < BLOCK_SIZE; ii++)  
        for (int kk = 0; kk < BLOCK_SIZE; kk++)  
            for (int jj = 0; jj < BLOCK_SIZE; jj += nW) {  
                __m512 vB = _mm512_load_ps(&bB[kk][jj]);  
                __m512 vA = _mm512_set1_ps(bA[ii][kk]);  
                __m512 vC = _mm512_load_ps(&bC[ii][jj]);  
                vC = _mm512_fmadd_ps(vA, vB, vC);  
                _mm512_store_ps(&bC[ii][jj], vC);  
            }  
}
```



# What are (assembly) intrinsics?

*An intrinsic function is a subroutine built-in to the compiler, that has a special meaning (beyond what C++ would suggest)*

*Assembly intrinsics in C, in particular, are subroutines that (in principle) encapsulate the operation of one (or a few) CPU assembly instruction*

*The compiler typically also provides special data types to encapsulate special types of registers (what's relevant to us: SIMD registers)*

*Major benefit of using intrinsics*

*(as opposed to in-line assembly, or editing the assembly code file directly) :*

- No need for allocating registers (compiler does it)*
- Even ok to use more “vector” variables than available on processor (the compiler will take care of stashing “spilling” them to temporary memory)*
- Significantly easier syntax*
- Intrinsics are available that map to several assembly instructions (or can be collapsed into even fewer ones)*

# What are (assembly) intrinsics?

*We will offer a brief walkthrough-by-example*

*Reference materials (if you want to dive deeper; not essential for this class)*

*Intel 64 & IA-32 Architectures Optimization Reference Manual*

*<https://www.intel.com/content/dam/www/public/us/en/documents/manuals/64-ia-32-architectures-optimization-manual.pdf>*

*Intel 64 and IA-32 Architectures Software Developer's Manual*

*<https://software.intel.com/en-us/articles/intel-sdm>*

*Intel Intrinsics Guide*

*<https://software.intel.com/sites/landingpage/IntrinsicsGuide/>*

# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

*DenseAlgebra/GEMM\_Test\_1\_0\_avx512*

```
#include "MatMatMultiplyBlockHelper.h"
#include "immintrin.h"

void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])
{
    static constexpr int nW = 16; // Width of SIMD vectors

    for (int ii = 0; ii < BLOCK_SIZE; ii++)
        for (int kk = 0; kk < BLOCK_SIZE; kk++)
            for (int jj = 0; jj < BLOCK_SIZE; jj += nW)
                __m512 vB = _mm512_load_ps(&bB[kk][jj]);
                __m512 vA = _mm512_set1_ps(bA[ii][kk]);
                __m512 vC = _mm512_load_ps(&bC[ii][jj]);
                vC = _mm512_fmadd_ps(vA, vB, vC);
                _mm512_store_ps(&bC[ii][jj], vC);
            }
}
```

*Data types (encapsulating registers)*  
\_\_m512 Register with 16 floats (512 bits)  
\_\_m256 Register with 8 floats (256 bits)  
\_\_m512d Register with 8 doubles (512 bits)  
\_\_m256i Register with 8 32-bit ints (512 bits)

Special data types to encapsulate contents to be  
hosted in a SIMD register

# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

*DenseAlgebra/GEMM\_Test\_1\_0\_avx512*

```
#include "MatMatMultiplyBlockHelper.h"
#include "immintrin.h"

void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])
{
    static constexpr int nW = 16; // Width of SIMD vectors

    for (int ii = 0; ii < BLOCK_SIZE; ii++)
        for (int kk = 0; kk < BLOCK_SIZE; kk++)
            for (int jj = 0; jj < BLOCK_SIZE; jj += nW) {
                __m512 vB = _mm512_load_ps(&bB[kk][jj]);
                __m512 vA = _mm512_set1_ps(bA[ii][kk]);
                __m512 vC = _mm512_load_ps(&bC[ii][jj]);
                vC = _mm512_fmadd_ps(vA, vB, vC);
                _mm512_store_ps(&bC[ii][jj], vC);
            }
}
```

*Load vector register*

*\_mm512\_load\_ps Load register with 16 floats from a 64-byte aligned memory address (AVX512)*

*\_mm256\_load\_ps Load register with 8 floats from a 32-byte aligned memory address (AVX2)*

*Move aligned packed scalar vmovaps is the corresponding assembly instruction*

# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

*DenseAlgebra/GEMM\_Test\_1\_0\_avx512*

```
#include "MatMatMultiplyBlockHelper.h"
#include "immintrin.h"

void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])
{
    static constexpr int nW = 16; // Width of SIMD vectors

    for (int ii = 0; ii < BLOCK_SIZE; ii++)
        for (int kk = 0; kk < BLOCK_SIZE; kk++)
            for (int jj = 0; jj < BLOCK_SIZE; jj += nW) {
                __m512 vB = _mm512_load_ps(&bB[kk][jj]);
                __m512 vA = _mm512_set1_ps(bA[ii][kk]);
                __m512 vC = _mm512_load_ps(&bC[ii][jj]);
                vC = _mm512_fmadd_ps(vA, vB, vC);
                _mm512_store_ps(&bC[ii][jj], vC);
            }
}
```

*Since we're working on 16-wide vectors, we are stepping by 16 each time*

## Technologies

- ☐ MMX
- ☐ SSE
- ☐ SSE2
- ☐ SSE3
- ☐ SSSE3
- ☐ SSE4.1
- ☐ SSE4.2
- ☐ AVX
- ☐ AVX2
- ☐ FMA
- ☐ AVX-512
- ☐ KNC
- ☐ SVML
- ☐ Other

## Categories

- ☐ Application-Targeted
- ☐ Arithmetic
- ☐ Bit Manipulation
- ☐ Cast
- ☐ Compare
- ☐ Convert
- ☐ Cryptography
- ☐ Elementary Math Functions

The Intel Intrinsic Guide is an interactive reference tool for Intel intrinsic instructions, which are C style functions that provide access to many Intel instructions - including Intel® SSE, AVX, AVX-512, and more - without the need to write assembly code. ✕

 ✕ ?

**\_\_m512** \_mm512\_load\_ps (void const\* mem\_addr)

vmovaps

### Synopsis

```
__m512 _mm512_load_ps (void const* mem_addr)
#include <immintrin.h>
Instruction: vmovaps zmm, m512
CPUID Flags: AVX512F for AVX-512, KNCNI for KNC
```

### Description

Load 512-bits (composed of 16 packed single-precision (32-bit) floating-point elements) from memory into **dst.mem\_addr** must be aligned on a 64-byte boundary or a general-protection exception may be generated.

### Operation

```
dst[511:0] := MEM[mem_addr+511:mem_addr]
dst[MAX:512] := 0
```

### Performance

Architecture	Latency	Throughput (CPI)
Icelake	8	0.5
Skylake	1	0.5



# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

*DenseAlgebra/GEMM\_Test\_1\_0\_avx512*

```
#include "MatMatMultiplyBlockHelper.h"
#include "immintrin.h"

void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])
{
    static constexpr int nW = 16; // Width of SIMD vectors

    for (int ii = 0; ii < BLOCK_SIZE; ii++)
        for (int kk = 0; kk < BLOCK_SIZE; kk++)
            for (int jj = 0; jj < BLOCK_SIZE; jj += nW) {
                __m512 vB = _mm512_load_ps(&bB[kk][jj]);
                __m512 vA = _mm512_set1_ps(bA[ii][kk]);
                __m512 vC = _mm512_load_ps(&bC[ii][jj]);
                vC = _mm512_fmadd_ps(vA, vB, vC);
                _mm512_store_ps(&bC[ii][jj], vC);
            }
}
```

*Store vector register*

`_mm512_store_ps` *Store register with 16 floats to a 64-byte aligned memory address (AVX512)*

# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

*DenseAlgebra/GEMM\_Test\_1\_0\_avx512*

```
#include "MatMatMultiplyBlockHelper.h"
#include "immintrin.h"

void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])
{
    static constexpr int nW = 16; // Width of SIMD vectors

    for (int ii = 0; ii < BLOCK_SIZE; ii++)
        for (int kk = 0; kk < BLOCK_SIZE; kk++)
            for (int jj = 0; jj < BLOCK_SIZE; jj += nW) {
                __m512 vB = _mm512_load_ps(&bB[kk][jj]);
                __m512 vA = _mm512_set1_ps(bA[ii][kk]);
                __m512 vC = _mm512_load_ps(&bC[ii][jj]);
                vC = _mm512_fmadd_ps(vA, vB, vC);
                _mm512_store_ps(&bC[ii][jj], vC);
            }
}
```

*Fill vector register with copies of a single value*  
`_mm512_set1_ps` Fill all entries of a 16-float register with a single value  
(could be a constant, or a memory location)

# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

*DenseAlgebra/GEMM\_Test\_1\_0\_avx512*

```
#include "MatMatMultiplyBlockHelper.h"
#include "immintrin.h"

void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])
{
    static constexpr int nW = 16; // Width of SIMD vectors

    for (int ii = 0; ii < BLOCK_SIZE; ii++)
        for (int kk = 0; kk < BLOCK_SIZE; kk++)
            for (int jj = 0; jj < BLOCK_SIZE; jj += nW) {
                __m512 vB = _mm512_load_ps(&bB[kk][jj]);
                __m512 vA = _mm512_set1_ps(bA[ii][kk]);
                __m512 vC = _mm512_load_ps(&bC[ii][jj]);
                vC = _mm512_fmadd_ps(vA, vB, vC);
                _mm512_store_ps(&bC[ii][jj], vC);
            }
}
```

*`_mm512_set1_ps` is an example of an intrinsic without a “clear” 1-to-1 correspondence to a unique assembly instruction*

- If the argument is a constant, the compiler will take care of allocating/loading it*
- If argument is a memory location, the `vbroadcastss` instruction may be issued*
- In some cases the operation can be “embedded” in arithmetic assembly instructions*

# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

*DenseAlgebra/GEMM\_Test\_1\_0\_avx512*

```
#include "MatMatMultiplyBlockHelper.h"
#include "immintrin.h"

void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])
{
    static constexpr int nW = 16; // Width of SIMD vectors

    for (int ii = 0; ii < BLOCK_SIZE; ii++)
        for (int kk = 0; kk < BLOCK_SIZE; kk++)
            for (int jj = 0; jj < BLOCK_SIZE; jj += nW) {
                __m512 vB = _mm512_load_ps(&bB[kk][jj]);
                __m512 vA = _mm512_set1_ps(bA[ii][kk]);
                __m512 vC = _mm512_load_ps(&bC[ii][jj]);
                vC = _mm512_fmadd_ps(vA, vB, vC);
                _mm512_store_ps(&bC[ii][jj], vC);
            }
}
```

*Fused multiply-add*

\_mm512\_fmadd\_ps Multiplies first and second argument, add the third argument to the product, and returns the result to a vector register (kind of like the saxpy function we saw previously)

# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

*DenseAlgebra/GEMM\_Test\_1\_0\_avx512*

```
#include "MatMatMultiplyBlockHelper.h"
#include "immintrin.h"

void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])
{
    static constexpr int nW = 16; // Width of SIMD vectors

    for (int ii = 0; ii < BLOCK_SIZE; ii++)
        for (int kk = 0; kk < BLOCK_SIZE; kk++)
            for (int jj = 0; jj < BLOCK_SIZE; jj += nW) {
                __m512 vB = _mm512_load_ps(&bB[kk][jj]);
                __m512 vA = _mm512_set1_ps(bA[ii][kk]);
                __m512 vC = _mm512_load_ps(&bC[ii][jj]);
                vC = _mm512_fmadd_ps(vA, vB, vC);
                _mm512_store_ps(&bC[ii][jj], vC);
            }
}
```

*Fused multiply-add*

*`_mm512_fmadd_ps` may actually be translated to one of many assembly instructions such as `vfmadd213ps`, `vfmadd132`, `vfmadd231` depending on which of the 3 inputs we are writing the result to (or multiple assembly instructions if we are writing to a different register)*

# Assembly code

MatMatMultiplyBlockHelper.s

*DenseAlgebra/GEMM\_Test\_1\_0\_avx512*

[...]

```
vmovups    192(%rax,%rdx), %zmm3    #14.41
xorl       %r8d, %r8d              #10.5
vmovups    128(%rax,%rdx), %zmm2    #14.41
vmovups    64(%rax,%rdx), %zmm1     #14.41
vmovups    (%rax,%rdx), %zmm0       #14.41
                                     # LOE rax rdx rbx rbp rsi rdi r8 r9 r10 r12 r13 r14 r15 cl zmm0
```

zmm1 zmm2 zmm3

```
..B1.3:                                     # Preds ..B1.3 ..B1.2
                                     # Execution count [4.10e+03]
vbroadcastss (%r9,%r10,4), %zmm4    #13.40
incq       %r10                    #10.5
vfmadd231ps (%r8,%rsi), %zmm4, %zmm0 #15.18
vfmadd231ps 64(%r8,%rsi), %zmm4, %zmm1 #15.18
vfmadd231ps 128(%r8,%rsi), %zmm4, %zmm2 #15.18
vfmadd231ps 192(%r8,%rsi), %zmm4, %zmm3 #15.18
addq       $256, %r8               #10.5
cmpq       $64, %r10              #10.5
jb         ..B1.3                  # Prob 98% #10.5
                                     # LOE rax rdx rbx rbp rsi rdi r8 r9 r10 r12 r13 r14 r15 cl zmm0
```

zmm1 zmm2 zmm3

```
..B1.4:                                     # Preds ..B1.3
                                     # Execution count [6.40e+01]
incb       %cl                    #9.5
vmovups    %zmm3, 192(%rax,%rdx)    #16.30
vmovups    %zmm2, 128(%rax,%rdx)    #16.30
vmovups    %zmm1, 64(%rax,%rdx)     #16.30
vmovups    %zmm0, (%rax,%rdx)       #16.30
```

[...]



# Assembly code

MatMatMultiplyBlockHelper.s

DenseAlgebra/GEMM\_Test\_1\_0\_avx512

[...]

vmovups	192(%rax,%rdx), %zmm3	#14.41
xorl	%r8d, %r8d	#10.5
vmovups	128(%rax,%rdx), %zmm2	#14.41
vmovups	64(%rax,%rdx), %zmm1	#14.41
vmovups	(%rax,%rdx), %zmm0	#14.41

# LOE rax rdx rbx rbp rsi rdi r8 r9 r10 r12 r13 r14 r15 cl zmm0

zmm1 zmm2 zmm3

..B1.3: # Preds ..B1.2 D1.2

# Execution

vbroadcastss (%r9,%r10,4), %zmm4

incq %r10

vfmadd231ps (%r8,%rsi), %zmm4, %zmm0

vfmadd231ps 64(%r8,%rsi), %zmm4, %zmm1

vfmadd231ps 128(%r8,%rsi), %zmm4, %zmm2

vfmadd231ps 192(%r8,%rsi), %zmm4, %zmm3 #15.18

addq \$256, %r8 #10.5

cmpq \$64, %r10 #10.5

jb ..B1.3 # Prob 98% #10.5

# LOE rax rdx rbx rbp rsi rdi r8 r9 r10 r12 r13 r14 r15 cl zmm0

zmm1 zmm2 zmm3

..B1.4: # Preds ..B1.3

# Execution count [6.40e+01]

incb %cl #9.5

vmovups %zmm3, 192(%rax,%rdx) #16.30

vmovups %zmm2, 128(%rax,%rdx) #16.30

vmovups %zmm1, 64(%rax,%rdx) #16.30

vmovups %zmm0, (%rax,%rdx) #16.30

[...]

*Read the 64 floats starting at bB[kk][0]  
using 4x 16-wide store (“move”) instructions  
(the compiler does a bit extra loop reordering,  
hence the split of the “move” instructions)*

# Assembly code

MatMatMultiplyBlockHelper.s

DenseAlgebra/GEMM\_Test\_1\_0\_avx512

[...]

```
vmovups    192(%rax,%rdx), %zmm3    #14.41
xorl       %r8d, %r8d              #10.5
vmovups    128(%rax,%rdx), %zmm2    #14.41
vmovups    64(%rax,%rdx), %zmm1     #14.41
vmovups    (%rax,%rdx), %zmm0
```

# LOE rax rdx rbp

zmm1 zmm2 zmm3

```
..B1.3:    # Preds ..B1.3 ..B1.2
           # Execution count [4.10e+03]
```

```
vbroadcastss (%r9,%r10,4), %zmm4
```

#13.40

```
incq      %r10
```

idempotent operation avoids memory latency

#10.5

```
vfmadd231ps (%r8,%rsi), %zmm4, %zmm0
```

#15.18

```
vfmadd231ps 64(%r8,%rsi), %zmm4, %zmm1
```

#15.18

```
vfmadd231ps 128(%r8,%rsi), %zmm4, %zmm2
```

#15.18

```
vfmadd231ps 192(%r8,%rsi), %zmm4, %zmm3
```

#15.18

```
addq      $256, %r8
```

#10.5

```
cmpq      $64, %r10
```

#10.5

```
jb        ..B1.3
```

# Prob 98%

#10.5

# LOE rax rdx rbx rbp rsi rdi r8 r9 r10 r12 r13 r14 r15 cl zmm0

zmm1 zmm2 zmm3

```
..B1.4:    # Preds ..B1.3
```

# Execution count [6.40e+01]

```
incb      %cl
```

#9.5

```
vmovups    %zmm3, 192(%rax,%rdx)
```

#16.30

```
vmovups    %zmm2, 128(%rax,%rdx)
```

#16.30

```
vmovups    %zmm1, 64(%rax,%rdx)
```

#16.30

```
vmovups    %zmm0, (%rax,%rdx)
```

#16.30

[...]

# Assembly code

MatMatMultiplyBlockHelper.s

DenseAlgebra/GEMM\_Test\_1\_0\_avx512

[...]

```
vmovups    192(%rax,%rdx), %zmm3          #14.41
xorl       %r8d, %r8d                     #10.5
vmovups    128(%rax,%rdx), %zmm2          #14.41
vmovups    64(%rax,%rdx), %zmm1           #14.41
vmovups    (%rax,%rdx), %zmm0             #14.41
# LOE rax rdx rbx rbp rsi rdi r8 r9 r10 r12 r13 r14 r15 cl zmm0
```

zmm1 zmm2 zmm3

```
..B1.3:                                     # Preds ..B1.3 ..B1.2
                                     # Execution count [4.10e+03]
```

```
vbroadcastss (%r9,%r10,4), %zmm4          #13.40
incq       %r10                          #10.5
```

```
vfmadd231ps (%r8,%rsi), %zmm4, %zmm0
vfmadd231ps 64(%r8,%rsi), %zmm4, %zmm1
vfmadd231ps 128(%r8,%rsi), %zmm4, %zmm2
vfmadd231ps 192(%r8,%rsi), %zmm4, %zmm3
```

```
addq       $256, %r8
cmpq       $64, %r10
jb         ..B1.3                         # Prob 98%
# LOE rax rdx rbx rbp rsi rdi r8 r9 r10 r12 r13 r14 r15 cl zmm0
```

zmm1 zmm2 zmm3

```
..B1.4:                                     # Preds ..B1.3
                                     # Execution count [6.40e+01]
```

```
incb      %cl                            #9.5
vmovups    %zmm3, 192(%rax,%rdx)          #16.30
vmovups    %zmm2, 128(%rax,%rdx)          #16.30
vmovups    %zmm1, 64(%rax,%rdx)           #16.30
vmovups    %zmm0, (%rax,%rdx)             #16.30
```

[...]

*Perform fused-multiply-add operation on 64-values (with 4 instructions). Values of  $bC[ii][jj]$  are directly read/written from/to memory*

# Assembly code

MatMatMultiplyBlockHelper.s

DenseAlgebra/GEMM\_Test\_1\_0\_avx512

[...]

```
vmovups    192(%rax,%rdx), %zmm3    #14.41
xorl       %r8d, %r8d              #10.5
vmovups    128(%rax,%rdx), %zmm2    #14.41
vmovups    64(%rax,%rdx), %zmm1     #14.41
vmovups    (%rax,%rdx), %zmm0       #14.41
# LOE rax rdx rbx rbp rsi rdi r8 r9 r10 r12 r13 r14 r15 cl zmm0
```

zmm1 zmm2 zmm3

```
..B1.3:    # Preds ..B1.3 ..B1.2
           # Execution count [4.10e+03]
vbroadcastss (%r9,%r10,4), %zmm4    #13.40
incq       %r10                    #10.5
vfmadd231ps (%r8,%rsi), %zmm4, %zmm0 #15.18
vfmadd231ps 64(%r8,%rsi), %zmm4, %zmm1 #15.18
vfmadd231ps 128(%r8,%rsi), %zmm4, %zmm2 #15.18
vfmadd231ps 192(%r8,%rsi), %zmm4, %zmm3 #15.18
addq       $256, %r8                #10.5
cmpq       $64, %r10               #10.5
jb         ..B1.3                  # Prob 98%
# LOE rax rdx rbx rbp rsi rdi r8 r9 r10 r12 r13 r14 r15 cl zmm0
```

zmm1 zmm2 zmm3

```
..B1.4:    # Preds ..B1.3
           # Execution count [6.40e+01]
incb       %cl                     #9.5
vmovups    %zmm3, 192(%rax,%rdx)    #16.30
vmovups    %zmm2, 128(%rax,%rdx)    #16.30
vmovups    %zmm1, 64(%rax,%rdx)     #16.30
vmovups    %zmm0, (%rax,%rdx)       #16.30
```

[...]

Overall : Slightly better code density,  
data/register reuse than what we had<sup>10</sup>  
before (with OpenMP)

# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

*DenseAlgebra/GEMM\_Test\_1\_0\_avx512*

```
#include "MatMatMultiplyBlockHelper.h"
#include "immintrin.h"

void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])
{
    static constexpr int nW = 16; // Width of SIMD vectors


    for (int ii = 0; ii < BLOCK_SIZE; ii++)
        for (int kk = 0; kk < BLOCK_SIZE; kk++)
            for (int jj = 0; jj < BLOCK_SIZE; jj += nW) {
                __m512 vB = _mm512_load_ps(&bB[kk][jj]);
                __m512 vA = _mm512_set1_ps(bA[ii][kk]);
                __m512 vC = _mm512_set1_ps(0.0f);
                _mm512_store_ps(&bC[ii][jj], vA * vB);
            }
}
```

*1.6x the runtime of MKL code!*


## Execution:

```
Running candidate kernel for correctness test ... [Elapsed time : 36.7904ms]
Running reference kernel for correctness test ... [Elapsed time : 40.292ms]
Discrepancy between two methods : 0.000154495
Running kernel for performance run # 1 ... [Elapsed time : 19.5309ms]
Running kernel for performance run # 2 ... [Elapsed time : 19.0874ms]
Running kernel for performance run # 3 ... [Elapsed time : 19.2922ms]
Running kernel for performance run # 4 ... [Elapsed time : 19.1488ms]
Running kernel for performance run # 5 ... [Elapsed time : 19.2003ms]
Running kernel for performance run # 6 ... [Elapsed time : 19.8611ms]
Running kernel for performance run # 7 ... [Elapsed time : 19.0866ms]
Running kernel for performance run # 8 ... [Elapsed time : 19.1242ms]
[...]
```

# Blocking for vectorization (once again ...)




# Blocking for vectorization (once again ...)







# Blocking for vectorization (once again ...)




# Blocking for vectorization (once again ...)




# Blocking for vectorization (once again ...)




# Blocking for vectorization (once again ...)




# Blocking for vectorization (once again ...)



# Blocking for vectorization (once again ...)

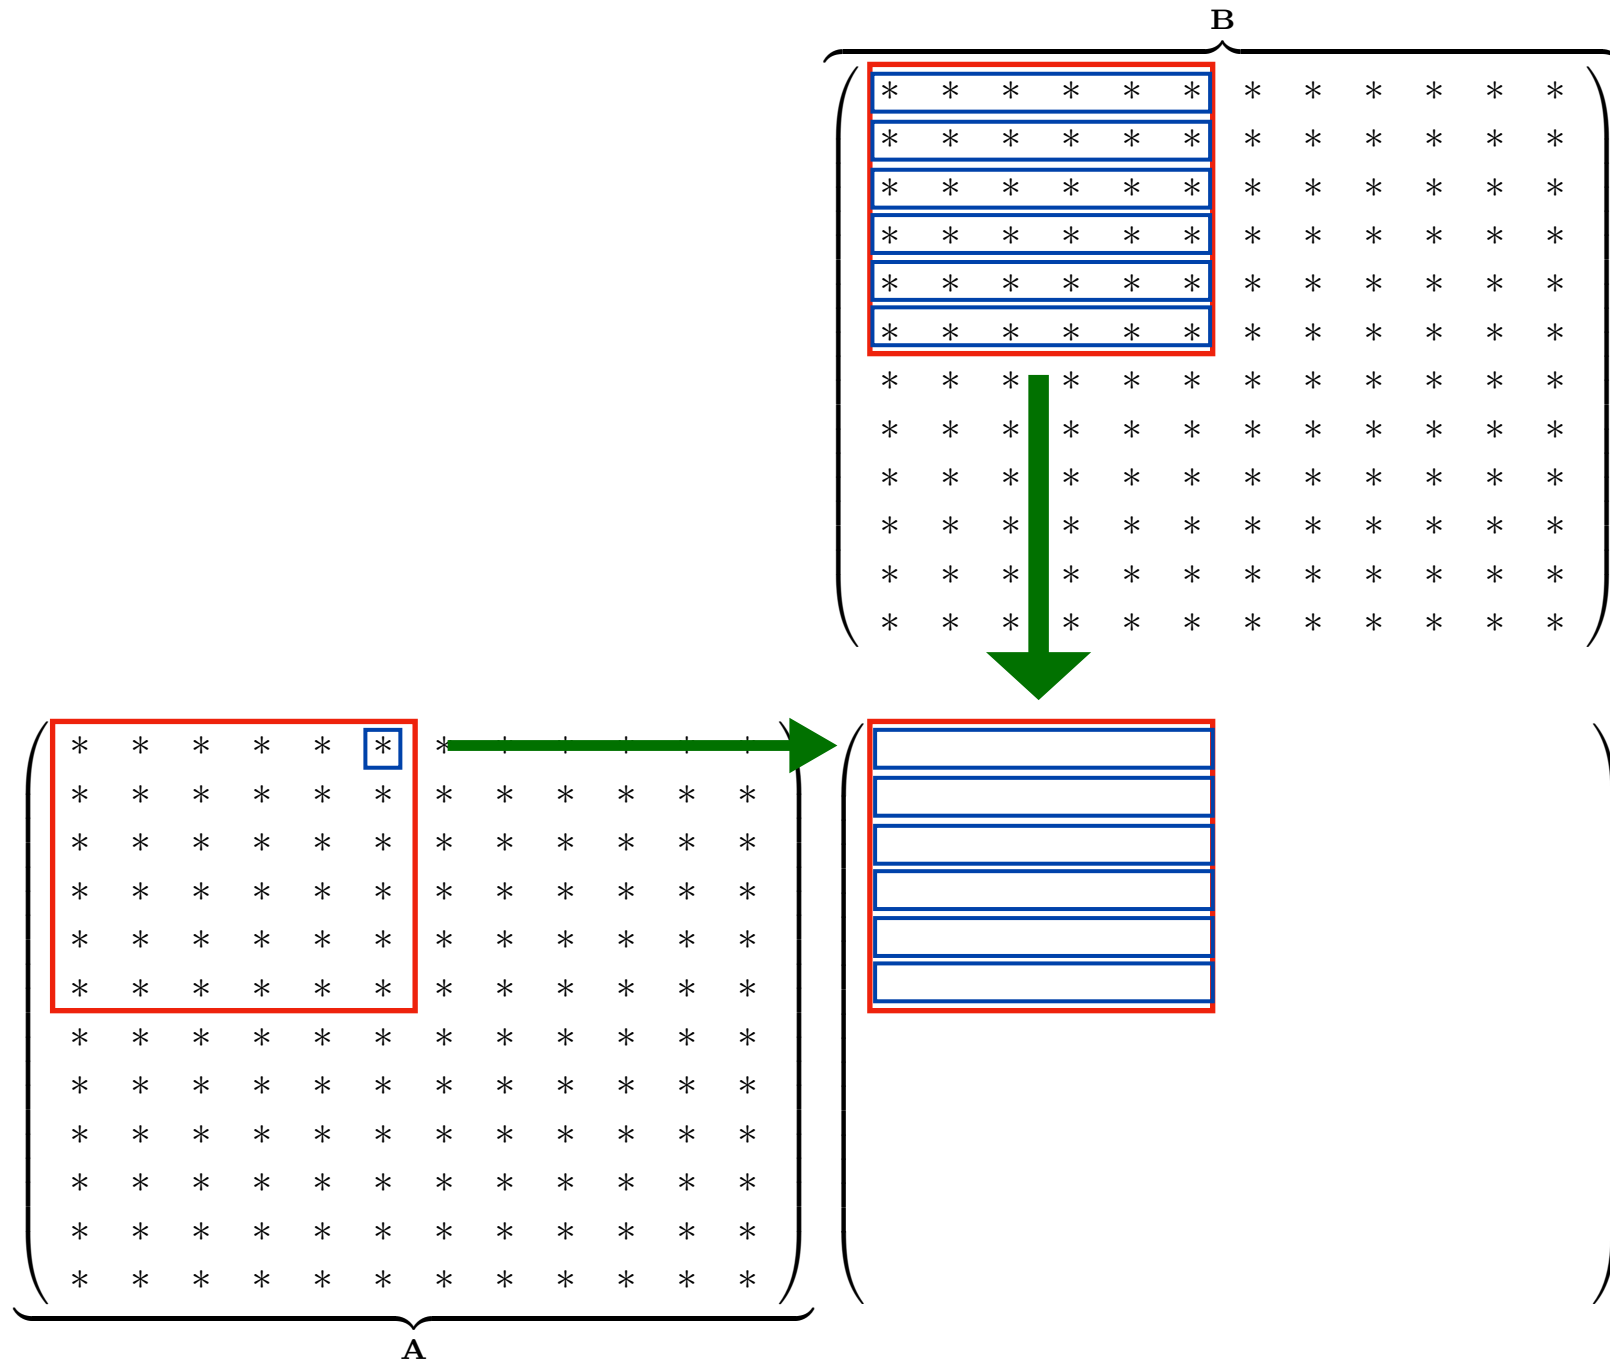


# Blocking for vectorization (once again ...)






# Blocking for vectorization (once again ...)



# Blocking for vectorization (once again ...)



# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

*DenseAlgebra/GEMM\_Test\_1\_1*

```
void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])
{
    static constexpr int nW = 16; // Entries in a SIMD vector (for AVX512; use 8 for AVX2)
    static constexpr int nB = BLOCK_SIZE / nW; // Number of blocks

    using const_blocked_matrix_t = const float (&) [nB][nW][nB][nW];
    using blocked_matrix_t = float (&) [nB][nW][nB][nW];

    // Matrix bA has indices [i][k], or in block form (bbA) [bi][ii][bk][kk]
    // Matrix bB has indices [k][j], or in block form (bbB) [bk][kk][bj][jj]
    // matrix bC has indices [i][j], or in block form (bbC) [bi][ii][bj][jj]

    auto bbA = reinterpret_cast<const_blocked_matrix_t>(bA[0][0]);
    auto bbB = reinterpret_cast<const_blocked_matrix_t>(bB[0][0]);
    auto bbC = reinterpret_cast<blocked_matrix_t>(bC[0][0]);

    for (int bi = 0; bi < nB; bi++)
        for (int bj = 0; bj < nB; bj++)
            for (int bk = 0; bk < nB; bk++)

            {
                for (int kk = 0; kk < nW; kk++)
                    for (int ii = 0; ii < nW; ii++)
#pragma omp simd
                        for (int jj = 0; jj < nW; jj++)
                            bbC[bi][ii][bj][jj] += bbA[bi][ii][bk][kk] * bbB[bk][kk][bj][jj];
            }
}
```

# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

*DenseAlgebra/GEMM\_Test\_1\_1*

```
void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])
{
    static constexpr int nW = 16; // Entries in a SIMD vector (for AVX512; use 8 for AVX2)
    static constexpr int nB = BLOCK_SIZE / nW; // Number of blocks

    using const_blocked_matrix_t = const float (&) [nB][nW][nB][nW];
    using blocked_matrix_t = float (&) [nB][nW][nB][nW];

    // Matrix bA has indices [i][k], or in block form (bbA) [bi][ii][bk][kk]
    // Matrix bB has indices [k][j], or in block form (bbB) [bk][kk][bj][jj]
    // matrix bC has indices [i][j], or in block form (bbC) [bi][ii][bj][jj]

    auto bbA = reinterpret_cast<const_blocked_matrix_t>(bA[0][0]);
    auto bbB = reinterpret_cast<const_blocked_matrix_t>(bB[0][0]);
    auto bbC = reinterpret_cast<blocked_matrix_t>(bC[0][0]);

    for (int bi = 0; bi < nB; bi++)
    for (int bj = 0; bj < nB; bj++)
    for (int bk = 0; bk < nB; bk++)

    {
        for (int kk = 0; kk < nW; kk++)
        for (int ii = 0; ii < nW; ii++)
#pragma omp simd
        for (int jj = 0; jj < nW; jj++)
            bbC[bi][ii][bj][jj] += bbA[bi][ii][bk][kk] * bbB[bk][kk][bj][jj];
    }
}
```

*Using blocking, once again ...  
(for the purposes of vectorization this time)*

# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

*DenseAlgebra/GEMM\_Test\_1\_1*

```
void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])
{
    static constexpr int nW = 16; // Entries in a SIMD vector (for AVX512; use 8 for AVX2)
    static constexpr int nB = BLOCK_SIZE / nW; // Number of blocks

    using const_blocked_matrix_t = const float (&) [nB][nW][nB][nW];
    using blocked_matrix_t = float (&) [nB][nW][nB][nW];

    // Matrix bA has indices [i][k], or in block form (bbA) [bi][ii][bk][kk]
    // Matrix bB has indices [k][j], or in block form (bbB) [bk][kk][bj][jj]
    // matrix bC has indices [i][j], or in block form (bbC) [bi][ii][bj][jj]

    auto bbA = reinterpret_cast<const_blocked_matrix_t>(bA[0][0]);
    auto bbB = reinterpret_cast<const_blocked_matrix_t>(bB[0][0]);
    auto bbC = reinterpret_cast<blocked_matrix_t>(bC[0][0]);

    for (int bi = 0; bi < nB; bi++)
    for (int bj = 0; bj < nB; bj++)
    for (int bk = 0; bk < nB; bk++)

    {
        for (int kk = 0; kk < nW; kk++)
        for (int ii = 0; ii < nW; ii++)
        #pragma omp simd
        for (int jj = 0; jj < nW; jj++)
            bbC[bi][ii][bj][jj] += bbA[bi][ii][bk][kk] * bbB[bk][kk][bj][jj];
    }
}
```

*Using blocking, once again ...  
(for the purposes of vectorization this time)*

# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

*DenseAlgebra/GEMM\_Test\_1\_1*

```
void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])
{
    static constexpr int nW = 16; // Entries in a SIMD vector (for AVX512; use 8 for AVX2)
    static constexpr int nB = BLOCK_SIZE / nW; // Number of blocks

    using const_blocked_matrix_t = const float (&) [nB][nW][nB][nW];
    using blocked_matrix_t = float (&) [nB][nW][nB][nW];

    // Matrix bA has indices [i][k], or in block form (bbA) [bi][ii][bk][kk]
    // Matrix bB has indices [k][j], or in block form (bbB) [bk][kk][bj][jj]
    // matrix bC has indices [i][j], or in block form (bbC) [bi][ii][bj][jj]

    auto bbA = reinterpret_cast<const_blocked_matrix_t>(bA[0]);
    auto bbB = reinterpret_cast<const_blocked_matrix_t>(bB[0][0]);
    auto bbC = reinterpret_cast<blocked_matrix_t>(bC[0][0]);

    Running candidate kernel for correctness test ... [Elapsed time : 35.129ms]
    for (int bi = 0; bi < nB; ++bi)
        Running reference kernel for correctness test ... [Elapsed time : 41.255ms]
        for (int bj = 0; bj < nB; ++bj)
            Discrepancy between two methods : 0.00015349
            for (int bk = 0; bk < nB; ++bk)
                Running kernel for performance run # 1 ... [Elapsed time : 18.9509ms]
                Running kernel for performance run # 2 ... [Elapsed time : 18.2873ms]
                {
                    Running kernel for performance run # 3 ... [Elapsed time : 18.3924ms]
                    for (int kk = 0; kk < nW; ++kk)
                        Running kernel for performance run # 4 ... [Elapsed time : 18.2958ms]
                        for (int ii = 0; ii < nW; ++ii)
                            Running kernel for performance run # 5 ... [Elapsed time : 18.8063ms]
                            #pragma omp simd
                            for (int jj = 0; jj < nW; ++jj)
                                Running kernel for performance run # 6 ... [Elapsed time : 18.8611ms]
                                bbC[bi][ii][bj][jj] += bbA[bi][ii][bk][kk] * bbB[bk][kk][bj][jj];
                                Running kernel for performance run # 7 ... [Elapsed time : 18.7826ms]
                                Running kernel for performance run # 8 ... [Elapsed time : 18.6212ms]
                                [...]
                            }
                        }
                    }
                }
            }
        }
    }
}
```

# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

*DenseAlgebra/GEMM\_Test\_1\_2\_avx512*

```
void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])
{
    static constexpr int nW = 16; // Width of a SIMD vector
    static constexpr int nB = BLOCK_SIZE/nW;
    using const_blocked_matrix_t = const float (&) [nB][nW][nB][nW];
    using blocked_matrix_t = float (&) [nB][nW][nB][nW];
    auto bbA = reinterpret_cast<const_blocked_matrix_t>(bA[0][0]);
    auto bbB = reinterpret_cast<const_blocked_matrix_t>(bB[0][0]);
    auto bbC = reinterpret_cast<blocked_matrix_t>(bC[0][0]);

    for (int bi = 0; bi < nB; bi++) for (int bj = 0; bj < nB; bj++)
    {
        __m512 vC[nW];
        for (int ii = 0; ii < nW; ii++)
            vC[ii] = _mm512_load_ps(&bbC[bi][ii][bj][0]);

        for (int bk = 0; bk < nB; bk++) {
            __m512 vB[nW];
            for (int kk = 0; kk < nW; kk++)
                vB[kk] = _mm512_load_ps(&bbB[bk][kk][bj][0]);

            for (int ii = 0; ii < nW; ii++) for (int kk = 0; kk < nW; kk++)
                vC[ii] = _mm512_fmadd_ps(_mm512_set1_ps(bbA[bi][ii][bk][kk]), vB[kk], vC[ii]); }

        for (int ii = 0; ii < nW; ii++)
            _mm512_store_ps(&bbC[bi][ii][bj][0], vC[ii]);
    }
}
```



# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

*DenseAlgebra/GEMM\_Test\_1\_2\_avx512*

```
void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])
{
    static constexpr int nW = 16; // Width of a SIMD vector
    static constexpr int nB = BLOCK_SIZE/nW;
    using const_blocked_matrix_t = const float (&) [nB][nW][nB][nW];
    using blocked_matrix_t = float (&) [nB][nW][nB][nW];
    auto bbA = reinterpret_cast<const_blocked_matrix_t>(bA[0][0]);
    auto bbB = reinterpret_cast<const_blocked_matrix_t>(bB[0][0]);
    auto bbC = reinterpret_cast<blocked_matrix_t>(bC[0][0]);

    for (int bi = 0; bi < nB; bi++) for (int bj = 0; bj < nB; bj++)
    {
        __m512 vC[nW];
        for (int ii = 0; ii < nW; ii++)
            vC[ii] = _mm512_load_ps(&bbC[bi][ii][bj][0]);


        for (int bk = 0; bk < nB; bk++) {
            __m512 vB[nW];
            for (int kk = 0; kk < nW; kk++)
                vB[kk] = _mm512_load_ps(&bbB[bk][kk][0][0]);

            for (int ii = 0; ii < nW; ii++) for (int kk = 0; kk < nW; kk++)
                vC[ii] = _mm512_fmadd_ps(_mm512_set1_ps(bbA[bi][ii][bk][kk]), vB[kk], vC[ii]); }

        for (int ii = 0; ii < nW; ii++)
            _mm512_store_ps(&bbC[bi][ii][bj][0], vC[ii]);
    }
}
```

*- Define 16 “registers” vC[0] through vC[15]  
which will hold the contents of the C block*

# Blocking for vectorization (once again ...)



# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

*DenseAlgebra/GEMM\_Test\_1\_2\_avx512*

```
void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])
{
```

```
    static constexpr int nW = 16; // Width of a SIMD vector
    static constexpr int nB = BLOCK_SIZE/nW;
    using const_blocked_matrix_t = const float (&) [nB][nW][nB][nW];
    using blocked_matrix_t = float (&) [nB][nW][nB][nW];
    auto bbA = reinterpret_cast<const_blocked_matrix_t>(bA[0][0]);
    auto bbB = reinterpret_cast<const_blocked_matrix_t>(bB[0][0]);
    auto bbC = reinterpret_cast<blocked_matrix_t>(bC[0][0]);
```

```
    for (int bi = 0; bi < nB; bi++) for (int bj = 0; bj < nB; bj++)
    {
```

```
        __m512 vC[nW];
        for (int ii = 0; ii < nW; ii++)
            vC[ii] = _mm512_load_ps(&bbC[bi][ii][bj][0]);
```

```
        for (int bk = 0; bk < nB; bk++) {
            __m512 vB[nW];
            for (int kk = 0; kk < nW; kk++)
                vB[kk] = _mm512_load_ps(&bbB[bk][kk][bi][0]);
```

```
            for (int ii = 0; ii < nW; ii++) for (int jj = 0; jj < nW; jj++)
                vC[ii] = _mm512_fmadd_ps(_mm512_store_ps(vB, vC[ii]), vB[kk], vC[ii]);
```

```
            for (int ii = 0; ii < nW; ii++)
                _mm512_store_ps(&bbC[bi][ii][bj][0], vC[ii]);
        }
```

- Define 16 “registers” `vC[0]` through `vC[15]` which will hold the contents of the **C** block

- Populate them with the previous values from the blocked matrix **bbC**

Note: We are doing this outside the “bk” loop! No need to re-read C every time

```
}
```

# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

[DenseAlgebra/GEMM\\_Test\\_1\\_2\\_avx512](#)

```
void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])
{
    static constexpr int nW = 16; // Width of a SIMD vector
    static constexpr int nB = BLOCK_SIZE/nW;
    using const_blocked_matrix_t = const float (&) [nB][nW][nB][nW];
    using blocked_matrix_t = float (&) [nB][nW][nB][nW];
    auto bbA = reinterpret_cast<const_blocked_matrix_t>(bA[0][0]);
    auto bbB = reinterpret_cast<const_blocked_matrix_t>(bB[0][0]);
    auto bbC = reinterpret_cast<blocked_matrix_t>(bC[0][0]);

    for (int bi = 0; bi < nB; bi++) for (int bj = 0; bj < nB; bj++) {
        __m512 vC[nW];
        for (int ii = 0; ii < nW; ii++)
            vC[ii] = _mm512_load_ps(&bbC[bi][ii][bj][0]);


        for (int bk = 0; bk < nB; bk++) {
            __m512 vB[nW];
            for (int kk = 0; kk < nW; kk++)
                vB[kk] = _mm512_load_ps(&bbB[bk][kk][bj][0]);

            for (int ii = 0; ii < nW; ii++) for (int kk = 0; kk < nW; kk++)
                vC[ii] = _mm512_fmadd_ps(_mm512_set1_ps(bbA[bi][ii][bk][kk]), vB[kk], vC[ii]);
        }

        for (int ii = 0; ii < nW; ii++)
            _mm512_store_ps(&bbC[bi][ii][bj][0], vC[ii]);
    }
}
```

*- Similarly, define 16 “registers” vB[0] through vB[15] which will hold the contents of the **B** block*  
*- Read their values just once, before iterating through the matrix A (the ii & kk loop)*

# Blocking for vectorization (once again ...)



# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

*DenseAlgebra/GEMM\_Test\_1\_0\_avx512*

```
#include "MatMatMultiplyBlockHelper.h"
#include "immintrin.h"

void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])
{
    static constexpr int nW = 16; // Width of SIMD vectors

    for (int ii = 0; ii < BLOCK_SIZE; ii++)
        for (int kk = 0; kk < BLOCK_SIZE; kk++)
            for (int jj = 0; jj < BLOCK_SIZE; jj += nW) {
                __m512 vB = _mm512_load_ps(&bB[kk][jj]);
                __m512 vA = _mm512_set1_ps(bA[ii][kk]);
                __m512 vC = _mm512_load_ps(&bC[ii][jj]);
                vC = _mm512_fmadd_ps(vA, vB, vC);
                _mm512_store_ps(&bC[ii][jj], vC);
            }
}
```

*Compare to prior version (B is read repeatedly)*

# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

*DenseAlgebra/GEMM\_Test\_1\_2\_avx512*

```
void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])
{
    static constexpr int nW = 16; // Width of a SIMD vector
    static constexpr int nB = BLOCK_SIZE/nW;
    using const_blocked_matrix_t = const float (&) [nB][nW][nB][nW];
    using blocked_matrix_t = float (&) [nB][nW][nB][nW];
    auto bbA = reinterpret_cast<const_blocked_matrix_t>(bA[0][0]);
    auto bbB = reinterpret_cast<const_blocked_matrix_t>(bB[0][0]);
    auto bbC = reinterpret_cast<blocked_matrix_t>(bC[0][0]);

    for (int bi = 0; bi < nB; bi++) for (int bj = 0; bj < nB; bj++)
    {
        __m512 vC[nW];
        for (int ii = 0; ii < nW; ii++)
            vC[ii] = _mm512_load_ps(&bbC[bi][ii][bj][0]);

        for (int bk = 0; bk < nB; bk++) {
            __m512 vB[nW];
            for (int kk = 0; kk < nW; kk++)
                vB[kk] = _mm512_load_ps(&bbB[bk][kk][bj][0]);

            for (int ii = 0; ii < nW; ii++) for (int kk = 0; kk < nW; kk++)
                vC[ii] = _mm512_fmadd_ps(_mm512_set1_ps(bbA[bi][ii][bk][kk]), vB[kk], vC[ii]); }

            for (int ii = 0; ii < nW; ii++)
                _mm512_store_ps(&bbC[bi][ii][bj][0], vC[ii]);
        }
    }
}
```

- Perform fused multiply-add operation on registers for **B & C**  
- Inline the “broadcast” operation for the corresponding entry of **A**



# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

*DenseAlgebra/GEMM\_Test\_1\_2\_avx512*

```
void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])
{
    static constexpr int nW = 16; // Width of a SIMD vector
    static constexpr int nB = BLOCK_SIZE/nW;
    using const_blocked_matrix_t = const float (&) [nB][nW][nB][nW];
    using blocked_matrix_t = float (&) [nB][nW][nB][nW];
    auto bbA = reinterpret_cast<const_blocked_matrix_t>(bA[0][0]);
    auto bbB = reinterpret_cast<const_blocked_matrix_t>(bB[0][0]);
    auto bbC = reinterpret_cast<blocked_matrix_t>(bC[0][0]);

    for (int bi = 0; bi < nB; bi++) for (int bj = 0; bj < nB; bj++)
    {
        __m512 vC[nW];
        for (int ii = 0; ii < nW; ii++)
            vC[ii] = _mm512_load_ps(&bbC[bi][ii][bj][0]);

        for (int bk = 0; bk < nB; bk++) {
            __m512 vB[nW];
            for (int kk = 0; kk < nW; kk++)
                vB[kk] = _mm512_load_ps(&bbB[bk][kk][bj][0]);

            for (int ii = 0; ii < nW; ii++) for (int kk = 0; kk < nW; kk++)
                vC[ii] = _mm512_fmadd_ps(_mm512_set1_ps(bbA[bi][ii][bk][kk]), vB[kk], vC[ii]); }

            for (int ii = 0; ii < nW; ii++)
                _mm512_store_ps(&bbC[bi][ii][bj][0], vC[ii]);
        }
    }
}
```

*Store the result back to **bbC** at the end of the loops for bk, ii and kk*

# Assembly code

MatMatMultiplyBlockHelper.s

*DenseAlgebra/GEMM\_Test\_1\_2\_avx512*

```
[...]
..B1.4:                                # Preds ..B1.6 ..B1.3
                                         # Execution count [6.40e+01]
    vmovups    (%r10,%r11), %zmm15      #30.42
    xorb       %r15b, %r15b             #33.13
    vmovups    256(%r10,%r11), %zmm14   #30.42
[... .]
    vmovups    3584(%r10,%r11), %zmm1   #30.42
    vmovups    3840(%r10,%r11), %zmm0   #30.42
    xorl       %r14d, %r14d            #33.13
    movq       %rdx, %r13              #34.26
..B1.5:                                # Preds ..B1.5 ..B1.4
                                         # Execution count [1.02e+03]
    vbroadcastss (%r13), %zmm16         #34.26
    incb       %r15b                   #33.13
    vfmadd213ps 64(%rsp,%r14), %zmm15, %zmm16 #34.26
    vfmadd231ps 4(%r13){1to16}, %zmm14, %zmm16 #34.26
    vfmadd231ps 8(%r13){1to16}, %zmm13, %zmm16 #34.26
    vfmadd231ps 12(%r13){1to16}, %zmm12, %zmm16 #34.26
    vfmadd231ps 16(%r13){1to16}, %zmm11, %zmm16 #34.26
    vfmadd231ps 20(%r13){1to16}, %zmm10, %zmm16 #34.26
    vfmadd231ps 24(%r13){1to16}, %zmm9, %zmm16 #34.26
    vfmadd231ps 28(%r13){1to16}, %zmm8, %zmm16 #34.26
    vfmadd231ps 32(%r13){1to16}, %zmm7, %zmm16 #34.26
    vfmadd231ps 36(%r13){1to16}, %zmm6, %zmm16 #34.26
    vfmadd231ps 40(%r13){1to16}, %zmm5, %zmm16 #34.26
    vfmadd231ps 44(%r13){1to16}, %zmm4, %zmm16 #34.26
    vfmadd231ps 48(%r13){1to16}, %zmm3, %zmm16 #34.26
    vfmadd231ps 52(%r13){1to16}, %zmm2, %zmm16 #34.26
    vfmadd231ps 56(%r13){1to16}, %zmm1, %zmm16 #34.26
    vfmadd231ps 60(%r13){1to16}, %zmm0, %zmm16 #34.26
    addq       $256, %r13               #33.13
    vmovups    %zmm16, 64(%rsp,%r14)    #34.17
```

[...]

# Assembly code

MatMatMultiplyBlockHelper.s

DenseAlgebra/GEMM\_Test\_1\_2\_avx512

```
[...]
..B1.4:                                # Preds ..B1.6 ..B1.3
                                         # Execution count [6.40e+01]
vmovups    (%r10,%r11), %zmm15          #30.42
xorb       %r15b, %r15b                 #33.13
vmovups    256(%r10,%r11), %zmm14       #30.42
[... ..]
vmovups    3584(%r10,%r11), %zmm1
vmovups    3840(%r10,%r11), %zmm0
xorl       %r14d, %r14d
movq       %rdx, %r13
..B1.5:                                # Preds ..B1.5 ..B1.4
                                         # Execution count [1.02e+03]
vbroadcastss (%r13), %zmm16             #34.26
incb       %r15b                       #33.13
vfmadd213ps 64(%rsp,%r14), %zmm15, %zmm16 #34.26
vfmadd231ps 4(%r13){1to16}, %zmm14, %zmm16 #34.26
vfmadd231ps 8(%r13){1to16}, %zmm13, %zmm16 #34.26
vfmadd231ps 12(%r13){1to16}, %zmm12, %zmm16 #34.26
vfmadd231ps 16(%r13){1to16}, %zmm11, %zmm16 #34.26
vfmadd231ps 20(%r13){1to16}, %zmm10, %zmm16 #34.26
vfmadd231ps 24(%r13){1to16}, %zmm9, %zmm16 #34.26
vfmadd231ps 28(%r13){1to16}, %zmm8, %zmm16 #34.26
vfmadd231ps 32(%r13){1to16}, %zmm7, %zmm16 #34.26
vfmadd231ps 36(%r13){1to16}, %zmm6, %zmm16 #34.26
vfmadd231ps 40(%r13){1to16}, %zmm5, %zmm16 #34.26
vfmadd231ps 44(%r13){1to16}, %zmm4, %zmm16 #34.26
vfmadd231ps 48(%r13){1to16}, %zmm3, %zmm16 #34.26
vfmadd231ps 52(%r13){1to16}, %zmm2, %zmm16 #34.26
vfmadd231ps 56(%r13){1to16}, %zmm1, %zmm16 #34.26
vfmadd231ps 60(%r13){1to16}, %zmm0, %zmm16 #34.26
addq       $256, %r13                  #33.13
vmovups    %zmm16, 64(%rsp,%r14)        #34.17
```

- All of B pre-loaded into registers  
(%zmm0 through %zmm15)

[...]

# Assembly code

MatMatMultiplyBlockHelper.s

DenseAlgebra/GEMM\_Test\_1\_2\_avx512

```
[...]
..B1.4:                                # Preds ..B1.6 ..B1.3
                                           # Execution count [6.40e+01]
vmovups    (%r10,%r11), %zmm15          #30.42
xorb       %r15b, %r15b                #33.13
vmovups    256(%r10,%r11), %zmm14       #30.42

[... .]
vmovups    3584(%r10,%r11), %zmm1       #30.42
vmovups    3840(%r10,%r11), %zmm0       #30.42
xorl       %r14d, %r14d                #33.13
movq       %rdx, %r13                  #34.26

..B1.5:                                # Preds ..B1.5 ..B1.4
                                           # Execution count [1.02e+03]
vbroadcastss (%r13), %zmm16             #34.26
incb       %r15b
vfmadd231ps 64(%rsp,%r14), %zmm15, %zmm16
vfmadd231ps 4(%r13){1to16}, %zmm14, %zmm16
vfmadd231ps 8(%r13){1to16}, %zmm13, %zmm16
vfmadd231ps 12(%r13){1to16}, %zmm12, %zmm16
vfmadd231ps 16(%r13){1to16}, %zmm11, %zmm16
vfmadd231ps 20(%r13){1to16}, %zmm10, %zmm16
vfmadd231ps 24(%r13){1to16}, %zmm9, %zmm16
vfmadd231ps 28(%r13){1to16}, %zmm8, %zmm16
vfmadd231ps 32(%r13){1to16}, %zmm7, %zmm16
vfmadd231ps 36(%r13){1to16}, %zmm6, %zmm16
vfmadd231ps 40(%r13){1to16}, %zmm5, %zmm16
vfmadd231ps 44(%r13){1to16}, %zmm4, %zmm16
vfmadd231ps 48(%r13){1to16}, %zmm3, %zmm16
vfmadd231ps 52(%r13){1to16}, %zmm2, %zmm16
vfmadd231ps 56(%r13){1to16}, %zmm1, %zmm16
vfmadd231ps 60(%r13){1to16}, %zmm0, %zmm16
addq       $256, %r13                   #33.13
vmovups    %zmm16, 64(%rsp,%r14)        #34.17

[...]
```

- Even higher density of fused multiply-adds
- Broadcast operation embedded into the arithmetic operation!
- Better absorption of load latency (B's have been loaded much earlier)

# Assembly code

MatMatMultiplyBlockHelper.s

*DenseAlgebra/GEMM\_Test\_1\_0\_avx512*

[...]

```
vmovups    192(%rax,%rdx), %zmm3    #14.41
xorl       %r8d, %r8d              #10.5
vmovups    128(%rax,%rdx), %zmm2    #14.41
vmovups    64(%rax,%rdx), %zmm1     #14.41
vmovups    (%rax,%rdx), %zmm0       #14.41
                                     # LOE rax rdx rbx rbp rsi rdi r8 r9 r10 r12 r13 r14 r15 cl zmm0
```

zmm1 zmm2 zmm3

```
..B1.3:                                     # Preds ..B1.3 ..B1.2
                                     # Execution count [4.10e+03]
vbroadcastss (%r9,%r10,4), %zmm4    #13.40
incq       %r10                    #10.5
vfmadd231ps (%r8,%rsi), %zmm4, %zmm0 #15.18
vfmadd231ps 64(%r8,%rsi), %zmm4, %zmm1 #15.18
vfmadd231ps 128(%r8,%rsi), %zmm4, %zmm2 #15.18
vfmadd231ps 192(%r8,%rsi), %zmm4, %zmm3 #15.18
addq       $256, %r8                #10.5
cmpq       $64, %r10                #10.5
jb         ..B1.3                   # Prob 98% #10.5
                                     # LOE rax rdx rbx rbp rsi rdi r8 r9 r10 r12 r13 r14 r15 cl zmm0
```

zmm1 zmm2 zmm3

```
..B1.4:                                     # Preds ..B1.3
                                     # Execution count [6.40e+01]
incb       %cl                      #9.5
vmovups    %zmm3, 192(%rax,%rdx)    #16.30
vmovups    %zmm2, 128(%rax,%rdx)    #16.30
vmovups    %zmm1, 64(%rax,%rdx)     #16.30
vmovups    %zmm0, (%rax,%rdx)       #16.30
```

[...]



# Inner multiplication (MatMatMultiplyBlockHelper.cpp)

*DenseAlgebra/GEMM\_Test\_1\_2\_avx512*

```
void MatMatMultiplyBlockHelper(const float (&bA)[BLOCK_SIZE][BLOCK_SIZE],
    const float (&bB)[BLOCK_SIZE][BLOCK_SIZE], float (&bC)[BLOCK_SIZE][BLOCK_SIZE])
{
```

```
    static constexpr int nW = 16; // Width of a SIMD vector
    static constexpr int nB = BLOCK_SIZE/nW;
    using const_blocked_matrix_t = const float (&) [nB][nW][nB][nW];
    using blocked_matrix_t = float (&) [nB][nW][nB][nW];
    auto bbA = reinterpret_cast<const_blocked_matrix_t>(bA[0][0]);
    auto bbB = reinterpret_cast<const_blocked_matrix_t>(bB[0][0]);
    auto bbC = reinterpret_cast<blocked_matrix_t>(bC[0][0]);
```

```
    for (int bi = 0; bi < nB; bi++) for (int bj = 0; bj < nB; bj++)
    {
```

```
        __m512 vC[nW];
```

```
        for (int ii = 0; ii < nW; ii++)
```

```
            vC[ii] =
```

## Execution:

```
                Running candidate kernel for correctness test ... [Elapsed time : 36.8076ms]
```

```
            for (int bk = 0; bk < nW; bk++)
                Running reference kernel for correctness test ... [Elapsed time : 40.5675ms]
```

```
                __m512 vDiscrepancy = 0;
                Discrepancy between two methods : 0.000156403
```

```
            for (int ki = 0; ki < nW; ki++)
                Running kernel for performance run # 1 ... [Elapsed time : 19.7365ms]
```

```
                vB[kb][ki] =
```

```
                    Running kernel for performance run # 2 ... [Elapsed time : 17.6981ms]
```

```
                    Running kernel for performance run # 3 ... [Elapsed time : 16.658ms]
```

```
            for (int kj = 0; kj < nW; kj++)
                Running kernel for performance run # 4 ... [Elapsed time : 16.4186ms]
```

```
                vC[kb][ki] =
```

```
                    Running kernel for performance run # 5 ... [Elapsed time : 17.454ms]
```

```
                    Running kernel for performance run # 6 ... [Elapsed time : 17.6172ms]
```

```
            for (int ii = 0; ii < nW; ii++)
                Running kernel for performance run # 7 ... [Elapsed time : 16.8232ms]
```

```
                __mm512 vResidual = 0;
                Running kernel for performance run # 8 ... [Elapsed time : 17.4193ms]
```

```
                [...]
```

*1.35x the runtime of MKL code!*