

## *Lecture 24: Laplacian Stencil in NanoVDB/IndexGrid (code review and high-level walkthrough)*

*Thursday April 20th 2023*



# Logistics

- Programming Assignment #4 due next Monday
- Midterm grades are posted. HW1 and HW2 should be posted soon!
- Guest lectures by Prof. Matt Sinclair on Apr 25, 27 and May 2nd

## Today's lecture




- An example of the Laplacian Stencil benchmark, implemented using NanoVDB/IndexGrid (<https://www.openvdb.org/>)
- Will \*not\* be included in your exams (material very fresh this year, very much an experiment in what's the right level of exposure to subject you to)
- Primarily a code walk-through. Please refer to OpenVDB/NanoVDB documentation for details on class structure and API specifics. Today is mostly explanation by example.

# Sparse grids as quadtrees - a specific example



*The “standard” quad tree representation  
(with data stored at leaves)*



# Sparse grids as quadtrees - a specific example



*Another interpretation: Instead of the tree holding the actual data, it holds “Indices” into a separate linearized array storage. This concept (which we will call an Index-Grid) will be used a lot in our specific approach!*





# OpenVDB/NanoVDB



*The data structure we will use to address some of these issues is OpenVDB (or its variant NanoVDB). The key idea is that it's a hierarchical (tree-like) sparse storage structure, but instead of each cube being split into  $2 \times 2 \times 2$  smaller "child" cubes, it's being split into  $8 \times 8$  (in 2D) or  $8 \times 8 \times 8$  (in 3D) children at every level of the tree! I.e. each tree node has 512 children!*


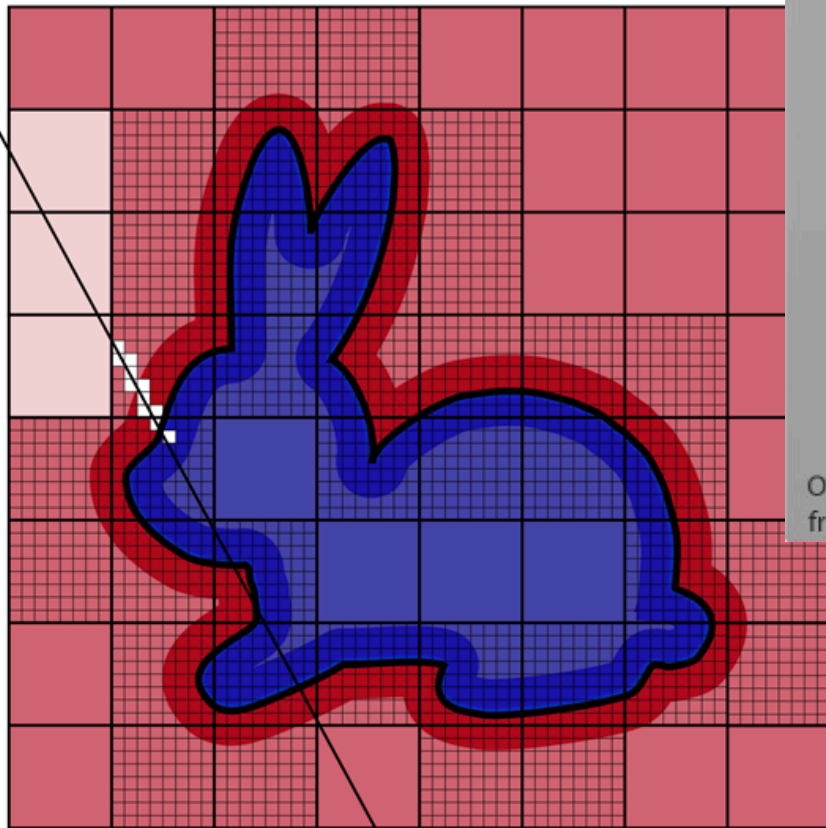
# OpenVDB/NanoVDB



*Benefit: At each leaf node of the VDB tree, we don't just have a single pixel (or 4/8 children), but a reasonably-sized 8x8x8 subgrid! This gives us plenty of opportunity to do most stencil operations within that grid!*



# OpenVDB/NanoVDB



OpenVDB: CPU-only, rich ecosystem, for dynamic sparsity, zoom into certain regions

nanoVDB: CPU/GPU Single header file, static sparsity

*Optimization #2: VDB has caching structures that allow quick (or rather quick) access to nearby neighbors of a leaf-level grid.*



# main.cpp

```
#include "Timer.h"  
#include "Laplacian.h"
```

```
#include <iomanip>
```

```
int main(int argc, char *argv[])  
{
```

```
    using array_t = float (&) [XDIM][YDIM][ZDIM];
```

```
    float *uRaw = new float [XDIM*YDIM*ZDIM];
```

```
    float *LuRaw = new float [XDIM*YDIM*ZDIM];
```

```
    float *uVDB = new float [XDIM*YDIM*ZDIM];
```

```
    float *LuVDB = new float [XDIM*YDIM*ZDIM];
```

```
    uint32_t *flagsVDB = new uint32_t [XDIM*YDIM*ZDIM];
```

```
    array_t u = reinterpret_cast<array_t>(*uRaw);
```

```
    array_t Lu = reinterpret_cast<array_t>(*LuRaw);
```

```
    Timer timer;
```

```
    timer.Start();
```

```
    auto handle = initializeIndexGrid();
```

```
    timer.Stop("Initializing indexGrid - Elapsed time :");
```

```
    auto *indexGridPtr = handle.grid<nanovdb::ValueIndex>();
```

```
    timer.Start();
```

```
    initializeData(u, Lu, indexGridPtr, uVDB, LuVDB, flagsVDB);
```

```
    timer.Stop("Initializing data - Elapsed time :");
```

```
    std::cout << "Discrepancy between dense and VDB (u) = " << compareData(u, indexGridPtr, uVDB)
```

*LaplacianNanoVDB/LaplacianNanoVDB\_0\_0*

*Make sure to compile with options:  
g++ -l. -O3 -fopenmp \*.cpp*

# main.cpp

*LaplacianNanoVDB/LaplacianNanoVDB\_0\_0*

```
#include "Timer.h"
#include "Laplacian.h"
```

```
#include <iomanip>
```

*New arrays for storing (linearized) sparse data*

```
int main(int argc, char *argv[])
{
```

```
    using array_t = float (&) [XDIM][YDIM][ZDIM];
```

```
    float *uRaw = new float [XDIM*YDIM*ZDIM];
```

```
    float *LuRaw = new float [XDIM*YDIM*ZDIM];
```

```
    float *uVDB = new float [XDIM*YDIM*ZDIM];
```

```
    float *LuVDB = new float [XDIM*YDIM*ZDIM];
```

```
    uint32_t *flagsVDB = new uint32_t [XDIM*YDIM*ZDIM];
```

```
    array_t u = reinterpret_cast<array_t>(*uRaw);
```

```
    array_t Lu = reinterpret_cast<array_t>(*LuRaw);
```

```
    Timer timer;
```

```
    timer.Start();
```

```
    auto handle = initializeIndexGrid();
```

```
    timer.Stop("Initializing indexGrid - Elapsed time :");
```

```
    auto *indexGridPtr = handle.grid<nanovdb::ValueIndex>();
```

```
    timer.Start();
```

```
    initializeData(u, Lu, indexGridPtr, uVDB, LuVDB, flagsVDB);
```

```
    timer.Stop("Initializing data - Elapsed time :");
```

```
    std::cout << "Discrepancy between dense and VDB (u) = " << compareData(u, indexGridPtr, uVDB)
```

# main.cpp

*LaplacianNanoVDB/LaplacianNanoVDB\_0\_0*

```
#include "Timer.h"
#include "Laplacian.h"
```

```
#include <iomanip>
```

```
int main(int argc, char *argv[])
{
```

```
    using array_t = float (&) [XDIM][YDIM][ZDIM];
```

```
    float *uRaw = new float [XDIM*YDIM*ZDIM];
```

```
    float *LuRaw = new float [XDIM*YDIM*ZDIM];
```

```
    float *uVDB = new float [XDIM*YDIM*ZDIM];
```

```
    float *LuVDB = new float [XDIM*YDIM*ZDIM];
```

```
    uint32_t *flagsVDB = new uint32_t [XDIM*YDIM*ZDIM];
```

```
    array_t u = reinterpret_cast<array_t>(*uRaw);
```

```
    array_t Lu = reinterpret_cast<array_t>(*LuRaw);
```

```
    Timer timer;
```

```
    timer.Start();
```

```
    auto handle = initializeIndexGrid();
```

```
    timer.Stop("Initializing indexGrid - Elapsed time :");
```

```
    auto *indexGridPtr = handle.grid<nanovdb::ValueIndex>();
```

```
    timer.Start();
```

```
    initializeData(u, Lu, indexGridPtr, uVDB, LuVDB, flagsVDB);
```

```
    timer.Stop("Initializing data - Elapsed time :");
```

```
    std::cout << "Discrepancy between dense and VDB (u) = " << compareData(u, indexGridPtr, uVDB)
```

*Initialization code for the appropriate NanoVDB  
(indexGrids) we will use*

# Laplacian.h

*LaplacianNanoVDB/LaplacianNanoVDB\_0\_0*

```
#pragma once
```

```
#include <nanovdb/NanoVDB.h>  
#include <nanovdb/util/GridHandle.h>
```

```
#define XDIM 512  
#define YDIM 512  
#define ZDIM 512  
#define INDEX_ACTIVE_FLAG 0xffffffff
```

```
void ComputeLaplacian(const float (&u)[XDIM][YDIM][ZDIM], float (&Lu)[XDIM][YDIM][ZDIM]);
```

```
nanovdb::GridHandle<nanovdb::HostBuffer> initializeIndexGrid();
```

```
void initializeData(float (&u)[XDIM][YDIM][ZDIM], float (&Lu)[XDIM][YDIM][ZDIM],  
    nanovdb::NanoGrid<nanovdb::ValueIndex>* indexGridPtr,  
    float *uBuffer, float *LuBuffer, uint32_t* flagsBuffer);
```

```
float compareData(const float (&data)[XDIM][YDIM][ZDIM],  
    nanovdb::NanoGrid<nanovdb::ValueIndex>* indexGridPtr,  
    const float *dataVDBBuffer);
```

```
void computeLaplacianVDB(nanovdb::NanoGrid<nanovdb::ValueIndex>* indexGridPtr,  
    float *uBuffer, float *LuBuffer, uint32_t* flagsBuffer);
```

*New additions in red; based on 3D dense  
Laplacian Stencil example  
(we'll go over them one by one)*



# main.cpp

*LaplacianNanoVDB/LaplacianNanoVDB\_0\_0*

```
#include "Timer.h"
#include "Laplacian.h"
```

```
#include <iomanip>
```

```
int main(int argc, char *argv[])
{
```

```
    using array_t = float (&) [XDIM][YDIM][ZDIM];
```

```
    float *uRaw = new float [XDIM*YDIM*ZDIM];
```

```
    float *LuRaw = new float [XDIM*YDIM*ZDIM];
```

```
    float *uVDB = new float [XDIM*YDIM*ZDIM];
```

```
    float *LuVDB = new float [XDIM*YDIM*ZDIM];
```

```
    uint32_t *flagsVDB = new uint32_t [XDIM*YDIM*ZDIM];
```

```
    array_t u = reinterpret_cast<array_t>(*uRaw);
```

```
    array_t Lu = reinterpret_cast<array_t>(*LuRaw);
```

```
    Timer timer;
```

```
    timer.Start();
```

```
    auto handle = initializeIndexGrid();
```

```
    timer.Stop("Initializing indexGrid - Elapsed time :");
```

```
    auto *indexGridPtr = handle.grid<nanovdb::ValueIndex>();
```

```
    timer.Start();
```

```
    initializeData(u, Lu, indexGridPtr, uVDB, LuVDB, flagsVDB);
```

```
    timer.Stop("Initializing data - Elapsed time :");
```

```
    std::cout << "Discrepancy between dense and VDB (u) = " << compareData(u, indexGridPtr, uVDB)
```

*Initialization code for the appropriate NanoVDB  
(indexGrids) we will use*

# Laplacian.cpp

```
#include "Laplacian.h"
```

```
#include <nanovdb/util/GridBuilder.h>
#include <nanovdb/util/IndexGridBuilder.h>
```

```
[...]
```

```
nanovdb::GridHandle<nanovdb::HostBuffer>
initializeIndexGrid()
```

```
{
    nanovdb::GridBuilder<nanovdb::ValueMask> builder(true);
    auto acc = builder.getAccessor();
    for (int i = 0; i < XDIM; i++)
    for (int j = 0; j < YDIM; j++)
    for (int k = 0; k < ZDIM; k++) {
        nanovdb::Coord xyz(i,j,k);
        acc.setValue(xyz, true);
    }
    auto handle = builder.getHandle();
    auto *dstGrid = handle.grid<nanovdb::ValueMask>();
    nanovdb::IndexGridBuilder<nanovdb::ValueMask> indexBuilder(*dstGrid, false, false);
    return indexBuilder.getHandle();
}
```

```
void initializeData(float (&u)[XDIM][YDIM][ZDIM], float (&Lu)[XDIM][YDIM][ZDIM],
    nanovdb::NanoGrid<nanovdb::ValueIndex>* indexGridPtr,
    float *uBuffer, float *LuBuffer, uint32_t* flagsBuffer)
{
```

*LaplacianNanoVDB/LaplacianNanoVDB\_0\_0*

*Initialize the sparsity pattern of array,  
store as IndexGrid  
(actually “dense” in this example)*

# main.cpp

*LaplacianNanoVDB/LaplacianNanoVDB\_0\_0*

*Put some initial values in u  
(dense and VDB) and compare*

```
#include "Timer.h"
#include "Laplacian.h"

#include <iomanip>

int main(int argc, char *argv[])
{
    [...]

    timer.Start();
    initializeData(u, Lu, indexGridPtr, uVDB, LuVDB, flagsVDB);
    timer.Stop("Initializing data - Elapsed time :");

    std::cout << "Discrepancy between dense and VDB (u) = "
        << compareData(u, indexGridPtr, uVDB) << std::endl;

    for(int test = 1; test <= 10; test++)
    {
        std::cout << "Running test iteration (dense) " << std::setw(2) << test << " ";
        timer.Start();
        ComputeLaplacian(u, Lu);
        timer.Stop("Elapsed time : ");
    }

    for(int test = 1; test <= 10; test++)
    {
        std::cout << "Running test iteration (VDB) " << std::setw(2) << test << " ";
        timer.Start();
        computeLaplacianVDB(indexGridPtr, uVDB, LuVDB, flagsVDB):
```

# Laplacian.cpp

*LaplacianNanoVDB/LaplacianNanoVDB\_0\_0*

```
#include "Laplacian.h"
```

```
[...]
```

```
void initializeData(float (&u)[XDIM][YDIM][ZDIM], float (&Lu)[XDIM][YDIM][ZDIM],  
    nanovdb::NanoGrid<nanovdb::ValueIndex>* indexGridPtr,  
    float *uBuffer, float *LuBuffer, uint32_t* flagsBuffer)
```

```
{
```

```
    float *uPtr = &u[0][0][0], *LuPtr = &Lu[0][0][0];
```

```
    // Zero out buffers by iterating linearly over them
```

```
    for (int n = 0; n < XDIM*YDIM*ZDIM; n++) {
```

```
        uPtr[n] = LuPtr[n] = uBuffer[n] = LuBuffer[n] = 0.0;
```

```
        flagsBuffer[n] = 0;
```

```
    }
```

```
    auto acc = indexGridPtr->getAccessor();
```

```
    for (int i = 1; i < XDIM-1; i++)
```

```
    for (int j = 1; j < YDIM-1; j++)
```

```
    for (int k = 1; k < ZDIM-1; k++) {
```

```
        nanovdb::Coord xyz(i,j,k);
```

```
        auto index = acc.getValue(xyz);
```

```
        u[i][j][k] = (float) ((i+j+k)%256-128);
```

```
        uBuffer[index] = u[i][j][k];
```

```
        flagsBuffer[index] = INDEX_ACTIVE_FLAG;
```

```
    }
```

```
}
```

```
float compareData(const float (&data)[XDIM][YDIM][ZDIM],
```

*Populating with initial data*



# Laplacian.cpp

*LaplacianNanoVDB/LaplacianNanoVDB\_0\_0*

```
#include "Laplacian.h"
```

```
[...]
```

```
float compareData(const float (&data)[XDIM][YDIM][ZDIM],
    nanovdb::NanoGrid<nanovdb::ValueIndex>* indexGridPtr,
    const float *dataVDBBuffer)
{
    float result = 0.;

    auto acc = indexGridPtr->getAccessor();
    for (int i = 0; i < XDIM; i++)
    for (int j = 0; j < YDIM; j++)
    for (int k = 0; k < ZDIM; k++) {
        nanovdb::Coord xyz(i,j,k);
        auto index = acc.getValue(xyz);
        result = std::max( result, std::abs(data[i][j][k]-dataVDBBuffer[index]) );
    }
}
```

*Comparing for correctness*

```
return result;
}

void computeLaplacianVDB(nanovdb::NanoGrid<nanovdb::ValueIndex>* indexGridPtr,
    float *uBuffer, float *LuBuffer, uint32_t* flagsBuffer)
{
    auto mgrHandle = createNodeManager(*indexGridPtr);
    auto *mgr = mgrHandle.template mgr<nanovdb::ValueIndex>();
    auto acc = indexGridPtr->getAccessor();
```

# main.cpp

```
#include "Timer.h"
#include "Laplacian.h"

#include <iomanip>

int main(int argc, char *argv[])
{
    [...]

    for(int test = 1; test <= 10; test++)
    {
        std::cout << "Running test iteration (dense) " << std::setw(2) << test << " ";
        timer.Start();
        ComputeLaplacian(u, Lu);
        timer.Stop("Elapsed time : ");
    }

    for(int test = 1; test <= 10; test++)
    {
        std::cout << "Running test iteration (VDB) " << std::setw(2) << test << " ";
        timer.Start();
        computeLaplacianVDB(indexGridPtr, uVDB, LuVDB, flagsVDB);
        timer.Stop("Elapsed time : ");
    }

    std::cout << "Discrepancy between dense and VDB (Lu) = "
        << compareData(Lu, indexGridPtr, LuVDB) << std::endl;
```

*LaplacianNanoVDB/LaplacianNanoVDB\_0\_0*

*Run and benchmark both versions,  
Check that the result is correct*

# Laplacian.cpp

*LaplacianNanoVDB/LaplacianNanoVDB\_0\_0*

[...]

```
void computeLaplacianVDB(nanovdb::NanoGrid<nanovdb::ValueIndex>* indexGridPtr,
    float *uBuffer, float *LuBuffer, uint32_t* flagsBuffer)
{
    auto mgrHandle = createNodeManager(*indexGridPtr);
    auto *mgr = mgrHandle.template mgr<nanovdb::ValueIndex>();
    auto acc = indexGridPtr->getAccessor();

#pragma omp parallel for firstprivate(acc)
    for ( size_t l = 0; l < mgr->nodeCount(0); ++l ) // l enumerates "leaves"
        for( auto iter = mgr->leaf(l).beginValue(); iter; ++iter ){
            auto coord = iter.getCoord(); // this is the coordinate within the leaf
            auto indexCtr = *iter; // this is the "center" index of the stencil;
            if (flagsBuffer[indexCtr] == INDEX_ACTIVE_FLAG) {
                auto indexPlusX  = acc.getValue(nanovdb::Coord(coord.x()+1, coord.y()  , coord.z()  ));
                auto indexMinusX = acc.getValue(nanovdb::Coord(coord.x()-1, coord.y()  , coord.z()  ));
                auto indexPlusY  = acc.getValue(nanovdb::Coord(coord.x()  , coord.y()+1, coord.z()  ));
                auto indexMinusY = acc.getValue(nanovdb::Coord(coord.x()  , coord.y()-1, coord.z()  ));
                auto indexPlusZ  = acc.getValue(nanovdb::Coord(coord.x()  , coord.y()  , coord.z()+1));
                auto indexMinusZ = acc.getValue(nanovdb::Coord(coord.x()  , coord.y()  , coord.z()-1));
                LuBuffer[indexCtr] =
                    -6 * uBuffer[indexCtr]
                    + uBuffer[indexPlusX]
                    + uBuffer[indexMinusX]
                    + uBuffer[indexPlusY]
                    + uBuffer[indexMinusY]
                    + uBuffer[indexPlusZ]
                    + uBuffer[indexMinusZ];
            }
        }
}
```

*VDB version of Laplacian Stencil!*

# Laplacian.cpp

LaplacianNanoVDB/LaplacianNanoVDB\_0\_0

$$[\dots]$$

```
void computeLaplacianVDB(nan
    float *uBuffer, float *L
{
```

```
auto mgrHandle = createN
auto *mgr = mgrHandle.te
auto acc = indexGridPtr-
```

```
#pragma omp parallel for fir
    for ( size_t l = 0; l <
```

```
for( auto iter = mgr
      auto coord = ite
```

```
auto indexCtr =
if (flagsBuffer[
```

auto indexPl  
auto indexMi

auto indexPl  
auto indexMi

auto indexPl  
auto indexMi

```
LuBuffer[ind
-6 * uBu
```

+ uBuffe  
+ uBuffe+ uBuffle  
+ uBuffle+ uBuffle  
+ uBuffle

```
[Initializing indexGrid - Elapsed time :244.07ms]
```

```
[Initializing data - Elapsed time :775.138ms]
```

Discrepancy between dense and VDB (u) = 0

Running test iteration (dense) 1 [Elapsed time : 67.4361ms]

Running test iteration (dense) 2 [Elapsed time : 67.0358ms]

Running test iteration (dense) 3 [Elapsed time : 67.3724ms]

```
Running test iteration (dense) 4 [Elapsed time : 67.0853ms]
```

Running test iteration (dense) 5 [Elapsed time : 68.2836ms]

Running test iteration (dense) 6 [Elapsed time : 67.192ms]

```
Running test iteration (dense) 7 [Elapsed time : 67.2768ms]
```

```
Running test iteration (dense) 8 [Elapsed time : 67.7054ms]
```

Running test iteration (dense) 9 [Elapsed time : 67.327ms]

```
Running test iteration (dense) 10 [Elapsed time : 67.0988ms]
```

Running test iteration (VDB) 1 [Elapsed time : 326.101ms]

Running test iteration (VDB) 2 [Elapsed time : 330.922ms]

Running test iteration (VDB) 3 [Elapsed time : 328.062ms]

Running test iteration (VDB) 4 [Elapsed time : 324.41ms]

```
Running test iteration (VDB) 5 [Elapsed time : 328.728ms]
```

Running test iteration (VDB) 6 [Elapsed time : 324.967ms]

Running test iteration (VDB) 7 [Elapsed time : 326.969ms]

```
Running test iteration (VDB) 8 [Elapsed time : 325.041ms]
```

Running test iteration (VDB) 9 [Elapsed time : 328.691ms]

Running test iteration (VDB) 10 [Elapsed time : 330.403ms]

Discrepancy between dense and VDB (Lu) = 0



# Optimality?

- Parallelization/multithreading of accessors?
- A different paradigm for operating on leaf nodes?  
(create local copy?)
- Ultimately: could get to 80-90% of dense performance